



Creating Databases and Tables



- We've focused on querying and reading data from existing databases and tables.
- Let's now shift our focus to creating our own databases and tables.



- Section Overview
 - Data Types
 - Primary and Foreign Keys
 - Constraints
 - CREATE
 - INSERT
 - UPDATE
 - DELETE, ALTER, DROP



- We first focus on learning a few theoretical concepts, such as choosing the correct data type for a stored value and setting possible constraints on it.
- We will also learn about primary and foreign keys.



Data Types



- We've already encountered a variety of data types, let's quickly review the main data types in SQL.



- Boolean
 - True or False
- Character
 - char, varchar, and text
- Numeric
 - integer and floating-point number
- Temporal
 - date, time, timestamp, and interval



- UUID
 - Universally Unique Identifiers
- Array
 - Stores an array of strings, numbers, etc.
- JSON
- Hstore key-value pair
- Special types such as network address and geometric data.



- When creating databases and tables, you should carefully consider which data types should be used for the data to be stored.
- Review the documentation to see limitations of data types:
- postgresql.org/docs/current/datatype.html



- For example
 - Imagine we want to store a phone number, should it be stored as numeric?
 - If so, which type of numeric?
- We could take a look at the documentation for options...



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



- Based on the limitations, you may think it makes sense to store it as a **BIGINT** data type, but we should really be thinking what is best for the situation.
- Why bother with numerics at all?
- We don't perform arithmetic with numbers, so it probably makes more sense as a **VARCHAR** data type instead.



- In fact, searching for best practice online, you will discover its usually recommended to store as a text based data type due to a variety of issues
 - No arithmetic performed
 - Leading zeros could cause issues, 7 and 07 treated same numerically, but are not the same phone number



- When creating a database and table, take your time to plan for long term storage
- Remember you can always remove historical information you've decided you aren't using, but you can't go back in time to add in information!



Primary and Foreign Keys



- A primary key is a column or a group of columns used to identify a row uniquely in a table.
- For example, in our dvdrental database we saw customers had a unique, non-null `customer_id` column as their primary key.



- Primary keys are also important since they allow us to easily discern what columns should be used for joining tables together.



- Example of Primary Key

Query Editor

Query History

1 SELECT * FROM customer

Data Output

Explain

Messages

Notifications

	customer_id [PK] integer	store_id smallint	first_name character varying (45)	last_name character varying (45)
1	524	1	Jared	Ely
2	1	1	Mary	Smith
3	2	1	Patricia	Johnson
4	3	1	Linda	Williams



- Example of Primary Key

Query Editor		Query History			
1		SELECT * FROM customer			
Data Output		Explain		Messages	
customer_id [PK] integer		store_id smallint	first_name character varying (45)	last_name character varying (45)	
1	524	1	Jared	Ely	
2	1	1	Mary	Smith	
3	2	1	Patricia	Johnson	
4	3	1	Linda	Williams	



- Notice its integer based and unique

Query Editor

Query History

1 SELECT * FROM customer

Data Output

Explain

Messages

Notifications

	customer_id [PK] integer	store_id smallint	first_name character varying (45)	last_name character varying (45)
1	524	1	Jared	Ely
2	1	1	Mary	Smith
3	2	1	Patricia	Johnson
4	3	1	Linda	Williams



- Later we will learn about SERIAL data type

Query Editor

Query History

1 SELECT * FROM customer

Data Output

Explain

Messages

Notifications

	customer_id [PK] integer	store_id smallint	first_name character varying (45)	last_name character varying (45)
1	524	1	Jared	Ely
2	1	1	Mary	Smith
3	2	1	Patricia	Johnson
4	3	1	Linda	Williams



- A foreign key is a field or group of fields in a table that uniquely identifies a row in another table.
- A foreign key is defined in a table that references to the primary key of the other table.



- The table that contains the foreign key is called referencing table or child table.
- The table to which the foreign key references is called referenced table or parent table.
- A table can have multiple foreign keys depending on its relationships with other tables.



- Recall in the dvdrental database payment table, each payment row had its unique payment_id (a primary key) and identified the customer that made the payment through the customer_id (a foreign key since it references the customer table's primary key)



- Example

Query Editor

Query History

1

SELECT * FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- Primary Key for Payment Table

Query Editor

Query History

1

SELECT * FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- Multiple Foreign Key References

Query Editor

Query History

1

SELECT * FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- Note pgAdmin won't alert you to FK

Query Editor

Query History

1 SELECT * FROM payment

Data Output

Explain

Messages

Notifications

	payment_id [PK] integer	customer_id smallint	staff_id smallint	rental_id integer	amount numeric (5,2)
1	17503	341	2	1520	7.99
2	17504	341	1	1778	1.99
3	17505	341	1	1849	7.99
4	17506	341	2	2829	2.99



- You may begin to realize primary key and foreign key typically make good column choices for joining together two or more tables.



- When creating tables and defining columns, we can use constraints to define columns as being a primary key, or attaching a foreign key relationship to another table.
- Let's quickly explore table properties in pgAdmin to see how to get information on primary and foreign keys!



Constraints



- Constraints are the rules enforced on data columns on table.
- These are used to prevent invalid data from being entered into the database.
- This ensures the accuracy and reliability of the data in the database.



- Constraints can be divided into two main categories:
 - Column Constraints
 - Constrains the data in a column to adhere to certain conditions.
 - Table Constraints
 - applied to the entire table rather than to an individual column.



- The most common constraints used:
 - **NOT NULL** Constraint
 - Ensures that a column cannot have NULL value.
 - **UNIQUE** Constraint
 - Ensures that all values in a column are different.



- The most common constraints used:
 - **PRIMARY Key**
 - Uniquely identifies each row/record in a database table.
 - **FOREIGN Key**
 - Constrains data based on columns in other tables.



- The most common constraints used:
 - **CHECK** Constraint
 - Ensures that all values in a column satisfy certain conditions.



- The most common constraints used:
 - **EXCLUSION** Constraint
 - Ensures that if any two rows are compared on the specified column or expression using the specified operator, not all of these comparisons will return TRUE.



- Table Constraints
 - CHECK (condition)
 - to check a condition when inserting or updating data.
 - REFERENCES
 - to constrain the value stored in the column that must exist in a column in another table.



- Table Constraints
 - UNIQUE (column_list)
 - forces the values stored in the columns listed inside the parentheses to be unique.
 - PRIMARY KEY(column_list)
 - Allows you to define the primary key that consists of multiple columns.



- Now that we understand data types, primary keys, foreign keys, and constraints we are ready to begin using SQL syntax to create tables!



CREATE



- Let's now learn the syntax to create a table in SQL using the CREATE keyword and column syntax.



- Full General Syntax
 - CREATE TABLE table_name (
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
 table_constraint table_constraint
) INHERITS existing_table_name;



- Full General Syntax
 - CREATE TABLE table_name (
column_name TYPE column_constraint,
column_name TYPE column_constraint,
table_constraint table_constraint
) INHERITS existing_table_name;



- Full General Syntax
 - CREATE TABLE table_name (
column_name TYPE column_constraint,
column_name TYPE column_constraint,
table_constraint table_constraint
) INHERITS existing_table_name;



- Full General Syntax
 - CREATE TABLE table_name (
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
 table_constraint table_constraint
) INHERITS existing_table_name;



- Full General Syntax
 - CREATE TABLE table_name (
column_name TYPE column_constraint,
column_name TYPE column_constraint,
table_constraint table_constraint
) INHERITS existing_table_name;



- Full General Syntax
 - CREATE TABLE table_name (
column_name TYPE column_constraint,
column_name TYPE column_constraint,
table_constraint table_constraint
) INHERITS existing_table_name;



- Full General Syntax

- CREATE TABLE table_name (
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
 table_constraint table_constraint
) INHERITS existing_table_name;



- Common Simple Syntax
 - CREATE TABLE table_name (
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE **table_name** (
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE **players**(
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE **players**(
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 column_name TYPE column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 player_id TYPE column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 player_id **TYPE** column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 player_id **SERIAL** column_constraint,
 column_name TYPE column_constraint,
);



- **SERIAL**

- In PostgreSQL, a sequence is a special kind of database object that generates a sequence of integers.
- A sequence is often used as the primary key column in a table.



- **SERIAL**

- It will create a sequence object and set the next value generated by the sequence as the default value for the column.
- This is perfect for a primary key, because it logs unique integer entries for you automatically upon insertion.



- **SERIAL**

- If a row is later removed, the column with the SERIAL data type will not adjust, marking the fact that a row was removed from the sequence, for example

- 1,2,3,5,6,7

- You know row 4 was removed at some point



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



- Example Syntax
 - CREATE TABLE players(
 player_id **SERIAL** column_constraint,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 player_id SERIAL **column_constraint**,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 player_id SERIAL PRIMARY KEY,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 player_id SERIAL PRIMARY KEY,
 column_name TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
player_id SERIAL PRIMARY KEY,
age TYPE column_constraint,
);



- Example Syntax
 - CREATE TABLE players(
 player_id SERIAL PRIMARY KEY,
 age TYPE column_constraint,
);



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



Name	Storage Size	Description	Range
<code>smallint</code>	2 bytes	small-range integer	-32768 to +32767
<code>integer</code>	4 bytes	typical choice for integer	-2147483648 to +2147483647
<code>bigint</code>	8 bytes	large-range integer	-9223372036854775808 to +9223372036854775807
<code>decimal</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>numeric</code>	variable	user-specified precision, exact	up to 131072 digits before the decimal point; up to 16383 digits after the decimal point
<code>real</code>	4 bytes	variable-precision, inexact	6 decimal digits precision
<code>double precision</code>	8 bytes	variable-precision, inexact	15 decimal digits precision
<code>smallserial</code>	2 bytes	small autoincrementing integer	1 to 32767
<code>serial</code>	4 bytes	autoincrementing integer	1 to 2147483647
<code>bigserial</code>	8 bytes	large autoincrementing integer	1 to 9223372036854775807



- Example Syntax
 - CREATE TABLE players(
 player_id SERIAL PRIMARY KEY,
 age **TYPE** column_constraint
);



- Example Syntax
 - CREATE TABLE players(
player_id SERIAL PRIMARY KEY,
age **SMALLINT** column_constraint
);



- Example Syntax
 - CREATE TABLE players(
player_id SERIAL PRIMARY KEY,
age SMALLINT **column_constraint**
);



- Example Syntax
 - CREATE TABLE players(
player_id SERIAL PRIMARY KEY,
age SMALLINT NOT NULL
);



- Example Syntax
 - CREATE TABLE players(
 player_id SERIAL PRIMARY KEY,
 age SMALLINT NOT NULL
);



- Let's explore some examples in pgAdmin!



INSERT



- INSERT allows you to add in rows to a table.
- General Syntax
 - INSERT INTO table (column1, column2, ...) VALUES
(value1, value2, ...),
(value1, value2, ...) ,...;



- INSERT allows you to add in rows to a table.
- Syntax for Inserting Values from another table:
 - INSERT INTO table(column1,column2,...)
SELECT column1,column2,...
FROM another_table
WHERE condition;



- Keep in mind, the inserted row values must match up for the table, including constraints.
- SERIAL columns do not need to be provided a value.
- Let's use INSERT in pgAdmin!



UPDATE



- The UPDATE keyword allows for the changing of values of the columns in a table.



- General Syntax
 - UPDATE table
SET column1 = value1,
column2 = value2 ,...
WHERE
condition;



- Example
 - UPDATE account
SET last_login = CURRENT_TIMESTAMP
WHERE last_login IS NULL;



- Reset everything without WHERE condition
 - UPDATE account
SET last_login = CURRENT_TIMESTAMP



- Set based on another column
 - UPDATE account
SET last_login = created_on



- Using another table's values (UPDATE join)
 - UPDATE TableA
SET original_col = TableB.new_col
FROM tableB
WHERE tableA.id = TableB.id



- Return affected rows
 - UPDATE account
SET last_login = created_on
RETURNING account_id,last_login



- Let's explore this further in pgAdmin!



DELETE



- We can use the DELETE clause to remove rows from a table.
- For example:
 - DELETE FROM table
WHERE row_id = 1



- We can delete rows based on their presence in other tables
- For example:
 - `DELETE FROM tableA
USING tableB
WHERE tableA.id=TableB.id`



- We can delete all rows from a table
- For example:
 - DELETE FROM table



- Similar to UPDATE command, you can also add in a RETURNING call to return rows that were removed.
- Let's explore DELETE with pgAdmin!



ALTER



- The ALTER clause allows for changes to an existing table structure, such as:
 - Adding, dropping, or renaming columns
 - Changing a column's data type
 - Set DEFAULT values for a column
 - Add CHECK constraints
 - Rename table



- General Syntax
 - ALTER TABLE table_name action



- Adding Columns
 - ALTER TABLE table_name
ADD COLUMN new_col TYPE



- Removing Columns
 - ALTER TABLE table_name
DROP COLUMN col_name



- Alter constraints
 - ALTER TABLE table_name
ALTER COLUMN col_name
SET DEFAULT value



- Alter constraints
 - ALTER TABLE table_name
 - ALTER COLUMN col_name
 - DROP DEFAULT



- Alter constraints
 - ALTER TABLE table_name
ALTER COLUMN col_name
SET NOT NULL



- Alter constraints
 - ALTER TABLE table_name
ALTER COLUMN col_name
DROP NOT NULL



- Alter constraints

- ALTER TABLE table_name

- ALTER COLUMN col_name

- ADD CONSTRAINT constraint_name



- Let's explore some examples in pgAdmin!



DROP



- DROP allows for the complete removal of a column in a table.
- In PostgreSQL this will also automatically remove all of its indexes and constraints involving the column.
- However, it will not remove columns used in views, triggers, or stored procedures without the additional CASCADE clause.



- General Syntax
 - ALTER TABLE table_name
 - DROP COLUMN col_name



- Remove all dependencies
 - ALTER TABLE table_name

DROP COLUMN col_name CASCADE



- Check for existence to avoid error
 - ALTER TABLE table_name

DROP COLUMN IF EXISTS col_name



- Drop multiple columns
 - ALTER TABLE table_name

DROP COLUMN col_one,

DROP COLUMN col_two



- Let's see a quick example in pgAdmin!



CHECK



- The CHECK constraint allows us to create more customized constraints that adhere to a certain condition.
- Such as making sure all inserted integer values fall below a certain threshold.



- General Syntax
 - CREATE TABLE example(
ex_id SERIAL PRIMARY KEY,
age SMALLINT CHECK (age > 21),
parent_age SMALLINT CHECK (
parent_age > age)
);



- Let's explore this concept in pgAdmin!